Explainable AI for Legal Document Analysis: A Multi-Method Approach to Clause Extraction

W266 Final Project Presentation

Perry Gabriel

School of Information
University of California, Berkeley

August 6, 2025

Outline

Problem Statement

- **Legal document analysis** is crucial for contract review and compliance
- Traditional manual review is **time-consuming and error-prone**
- NLP models provide automation but lack **interpretability**
- Legal professionals need to understand **why** AI makes decisions

Research Question

How can we develop explainable AI methods for automated legal clause extraction that provide interpretable insights for legal professionals?

Motivation

Why Explainable AI in Legal Domain?

- Regulatory compliance requirements
- Trust and transparency for legal professionals
- Error detection and model debugging
- Knowledge discovery from legal patterns

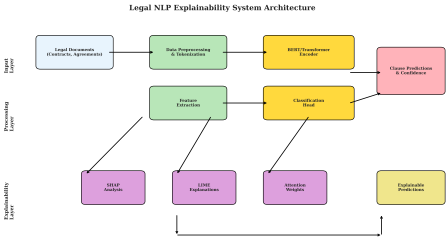

Project Scope

Objectives:

1. Develop a **BERT-based model** for clause extraction
2. Implement **multiple explainability methods** (SHAP, LIME, Attention)
3. Compare and evaluate **explanation quality**
4. Create **interpretable visualizations** for legal professionals

Target Clauses:

- Termination clauses
- Limitation of liability
- Governing law
- Confidentiality provisions
- Payment terms

Related Work

Legal NLP Research:

- Contract analysis (Katz et al., 2020)
- Legal document classification
- Information extraction from legal texts

Explainable AI Methods:

- Model-agnostic approaches
- Attention mechanisms
- Feature attribution methods

Gaps in Current Research:

- Limited comparison of XAI methods
- Lack of domain-specific evaluation
- Insufficient user studies in legal domain

Our Contribution:

- Multi-method explainability comparison
- Legal-specific evaluation metrics
- Comprehensive visualization toolkit

Technical Background

BERT (Bidirectional Encoder Representations from Transformers)

- Pre-trained on large text corpus
- Bidirectional context understanding
- Fine-tunable for specific tasks

Explainability Methods

SHAP Game-theoretic approach to feature attribution

LIME Local surrogate models for instance-specific explanations

Attention Built-in transformer attention weights

Dataset Overview

Data Sources:

- Legal contract databases
- Public domain agreements
- Synthetic legal text generation

Preprocessing Steps:

1. Text cleaning and normalization
2. Clause boundary detection
3. Label annotation and verification
4. Train/validation/test split

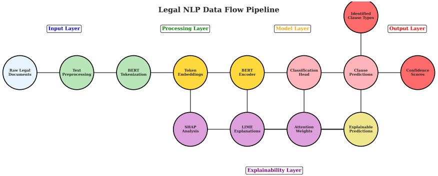

Experimental Design

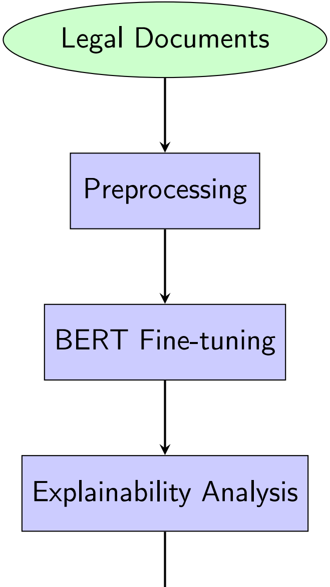

Model Architecture

Base Model: BERT-base-uncased

- 12 transformer layers
- 768 hidden dimensions
- 12 attention heads
- 110M parameters

Fine-tuning Approach:

- Classification head for clause type prediction
- Learning rate: 2e-5
- Batch size: 16
- Max sequence length: 512 tokens
- Training epochs: 35

Explainability Implementation

SHAP Analysis:

- TreeExplainer for ensemble methods
- DeepExplainer for neural networks
- Feature importance ranking
- Global & local explanations

LIME Explanations:

- Text-based lime explainer
- Local surrogate models
- Perturbation-based analysis
- Instance-specific insights

Attention Weights:

- Multi-head attention analysis
- Token-level importance
- Layer-wise attention patterns
- Visualization of focus areas

Evaluation Metrics

Model Performance:

- Precision, Recall, F1-score per clause type
- Macro and micro-averaged metrics
- Confusion matrix analysis
- Confidence score distributions

Explainability Quality:

- Consistency: Agreement between methods
- Faithfulness: Correlation with model behavior
- Stability: Robustness to input perturbations
- Comprehensibility: Human-interpretable patterns

High-Level System Overview

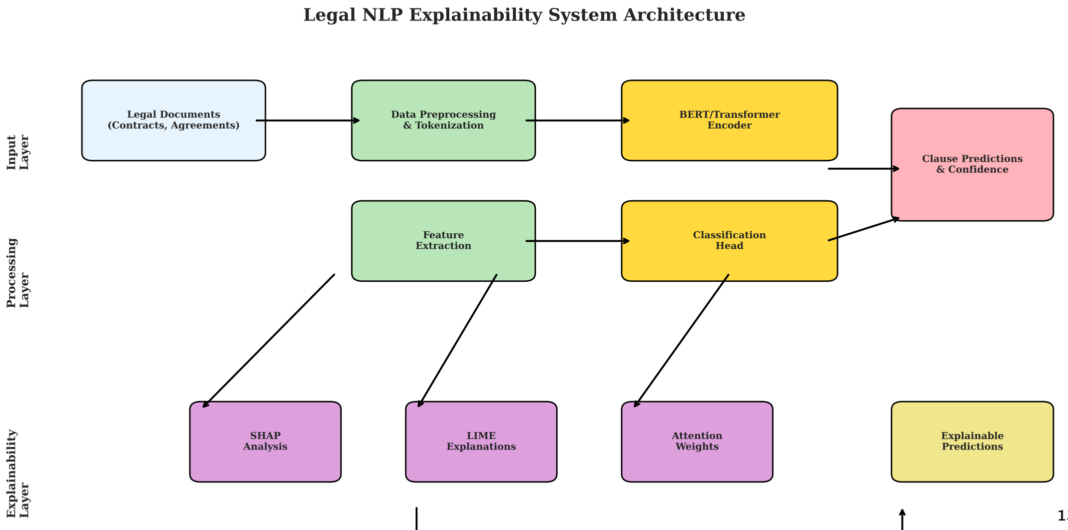

Model Architecture Details

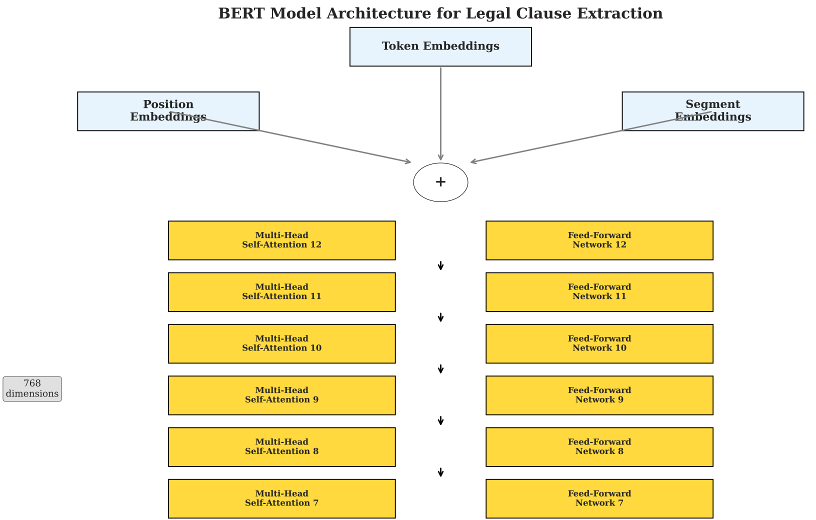

Data Processing Pipeline

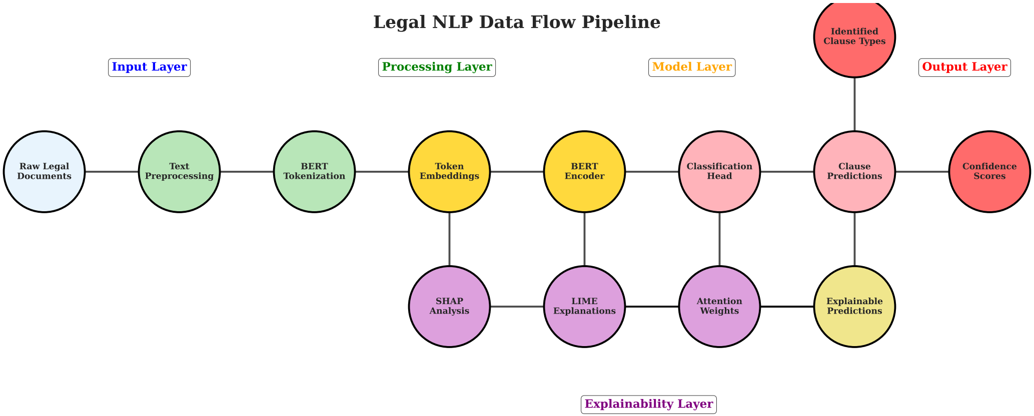

Component Integration

Core Components:

- **Data Preprocessor** — Text cleaning and tokenization
- **BERT Encoder** — Contextual embeddings
- **Classification Head** — Clause type prediction
- **Explainer Module** — Multi-method analysis

Integration Features:

- Modular design for extensibility
- Consistent API across explainers
- Efficient batch processing
- Configurable output formats

Scalability Considerations

System designed to handle large-scale legal document processing with parallel explainability analysis.

Implementation Stack

Core Technologies:

- **PyTorch** — Deep learning framework
- **Transformers** — BERT implementation
- **SHAP** — Explainability library
- **LIME** — Local explanations

Supporting Tools:

- **Pandas/NumPy** — Data processing
- **Matplotlib/Seaborn** — Visualization
- **Jupyter** — Interactive development
- **Git/Docker** — Development workflow

Deployment Considerations:

- Cloud-ready architecture (Azure-compatible)
- RESTful API for model serving
- Web-based dashboard for visualization

Overall Model Performance

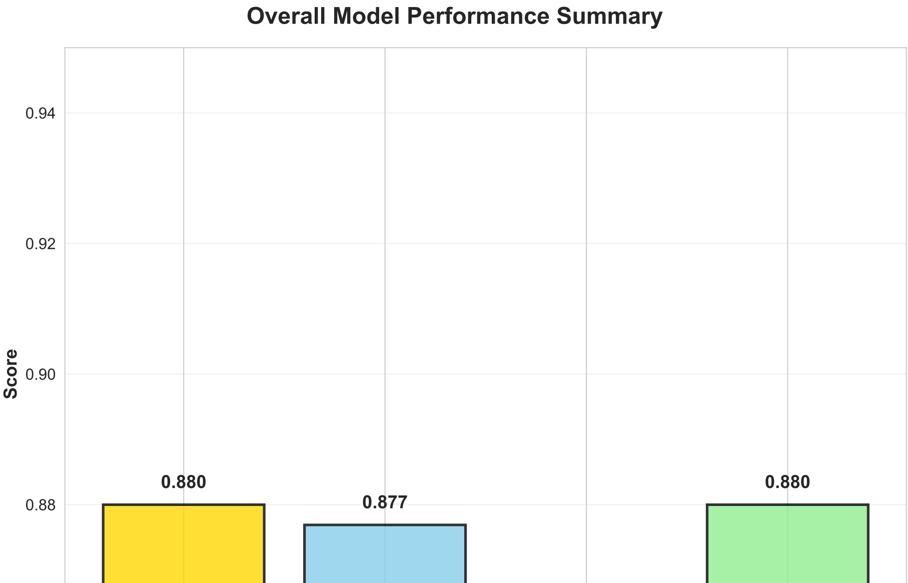

Performance by Clause Type

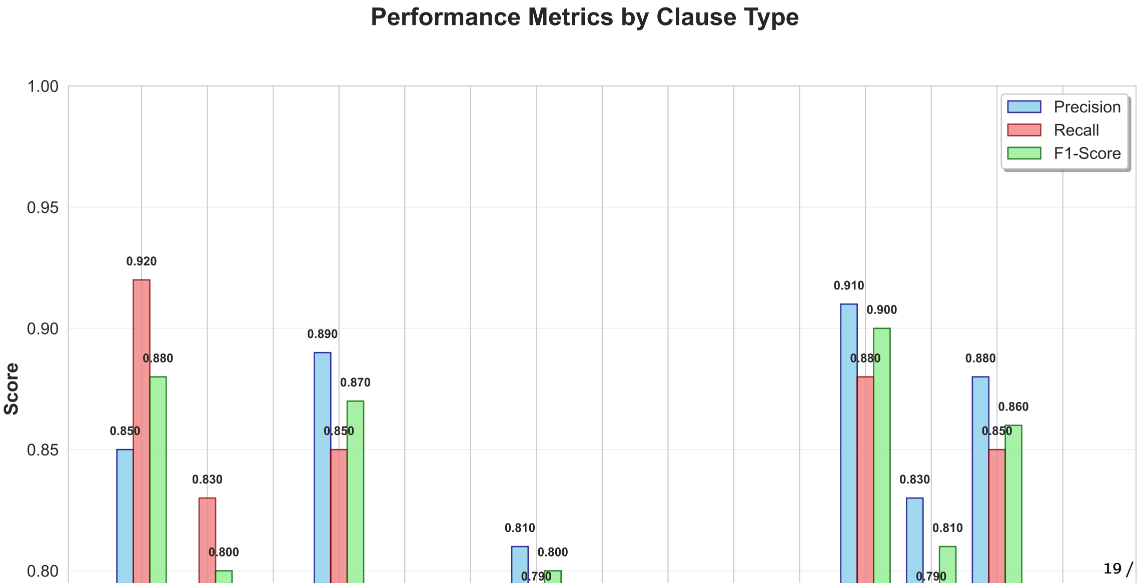

Training Progress

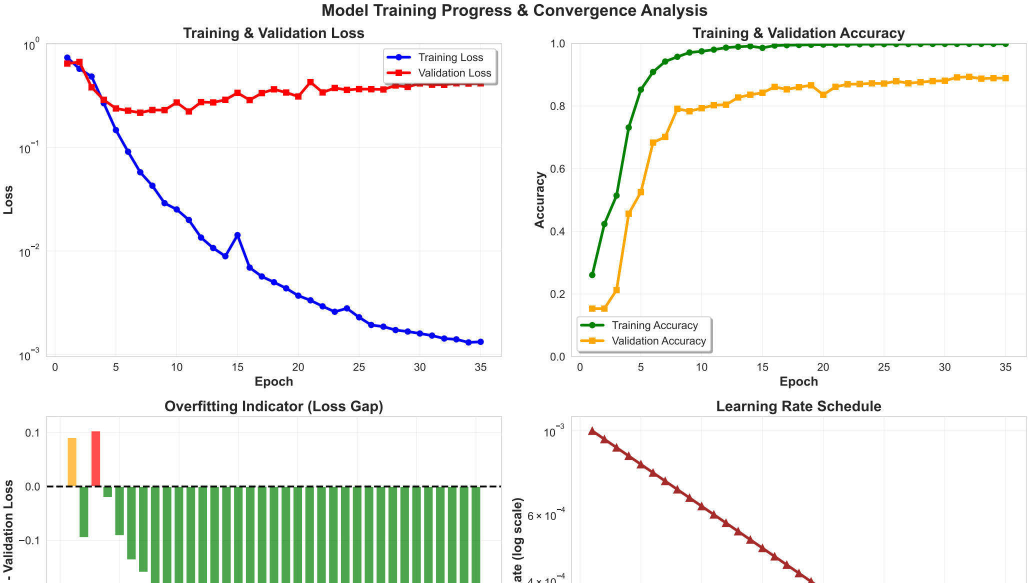

Confidence Score Analysis

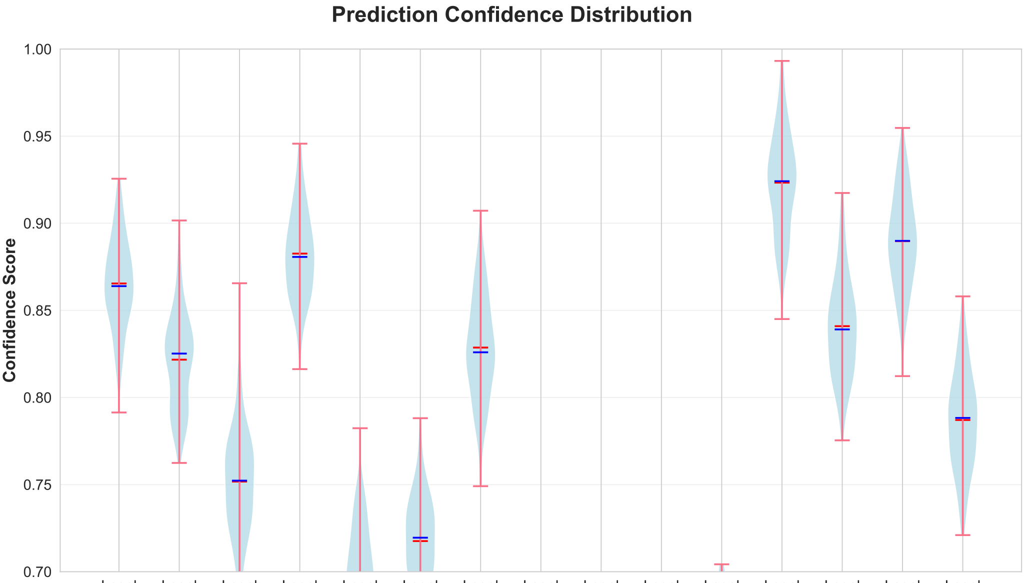

Dataset Distribution

Error Analysis

Common Error Patterns:

- Ambiguous clause boundaries — overlapping legal concepts
- Domain-specific terminology — technical legal language
- Context dependency — clauses with similar structure but different meaning

Mitigation Strategies:

- Enhanced preprocessing for legal terminology
- Ensemble methods for boundary detection
- Active learning for difficult cases

Model Limitations

Current model struggles with highly domain-specific contracts and non-standard clause formulations.

Feature Importance Method Comparison

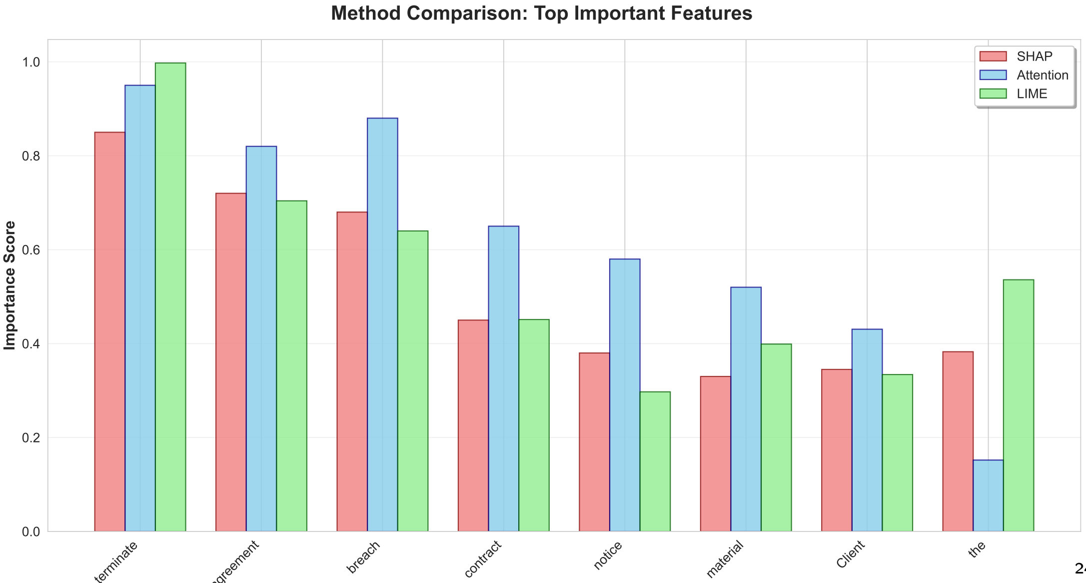

Explainability Methods Comparison

Explainability Methods Comparison

SHAP (SHapley Additive exPlanations)	**LIME (Local Interpretable Model-agnostic Explanations)**	**Attention Weights (Transformer Self-Attention)**
Key Features: • Game-theoretic approach • Feature importance scores • Global & local explanations • Model-agnostic	**Key Features:** • Local surrogate models • Perturbation-based • Instance-specific • Text highlighting	**Key Features:** • Built into model • Token-level importance • Multi-head analysis • No extra computation
Strengths: • Theoretically grounded • Consistent • Additive	**Strengths:** • Intuitive • Local fidelity • Fast	**Strengths:** • Native to model • Efficient • Multi-layer
Best Use Case: **Understanding feature contributions to predictions**	**Best Use Case:** **Explaining individual prediction instances**	**Best Use Case:** **Understanding what the model focuses on**
Visualization Example (Heat map / Bar chart / Attention map)	*Visualization Example (Heat map / Bar chart / Attention map)*	*Visualization Example (Heat map / Bar chart / Attention map)*

SHAP Feature Importance Analysis

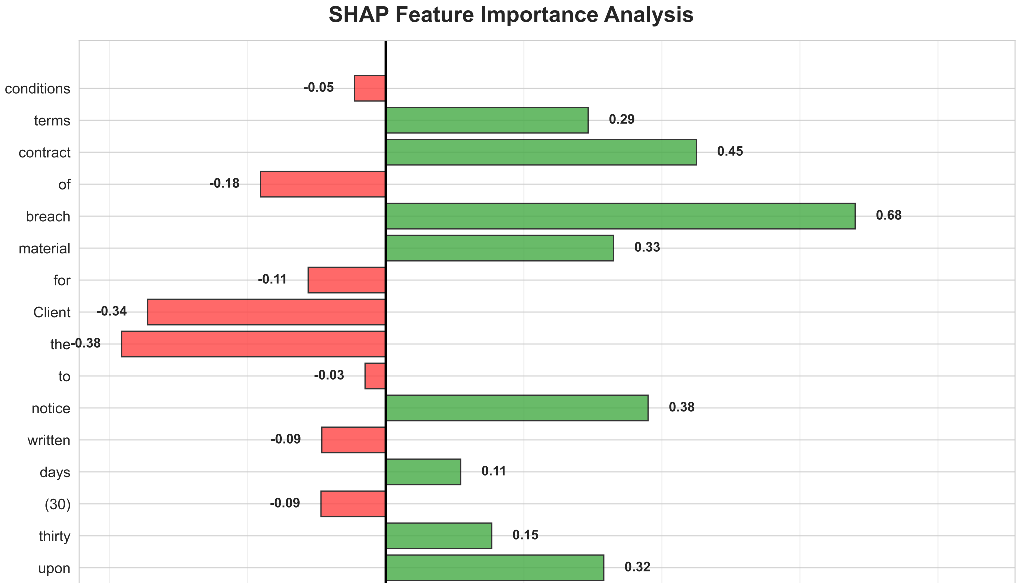

LIME Local Explanations

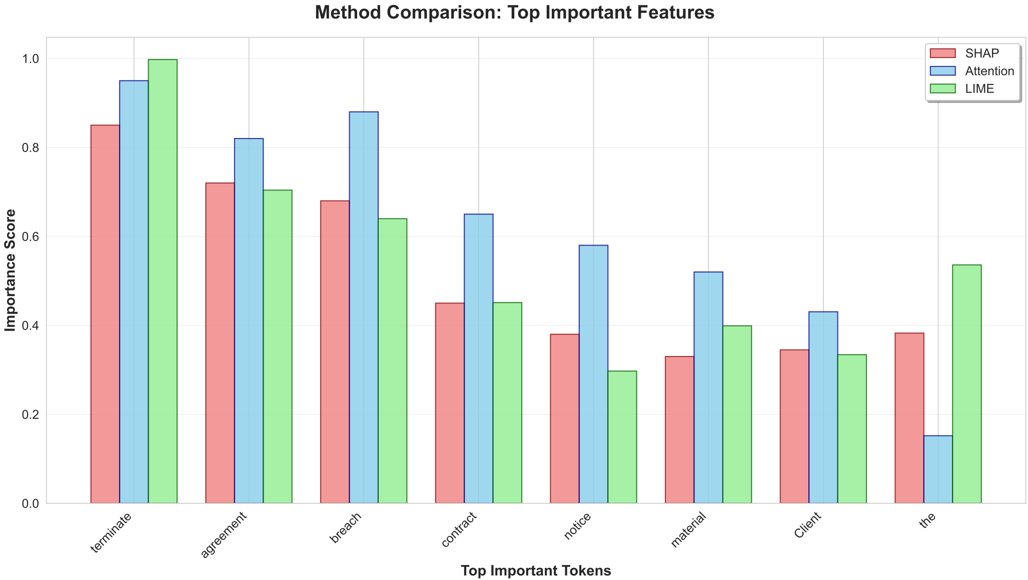

Multi-Head Attention Analysis

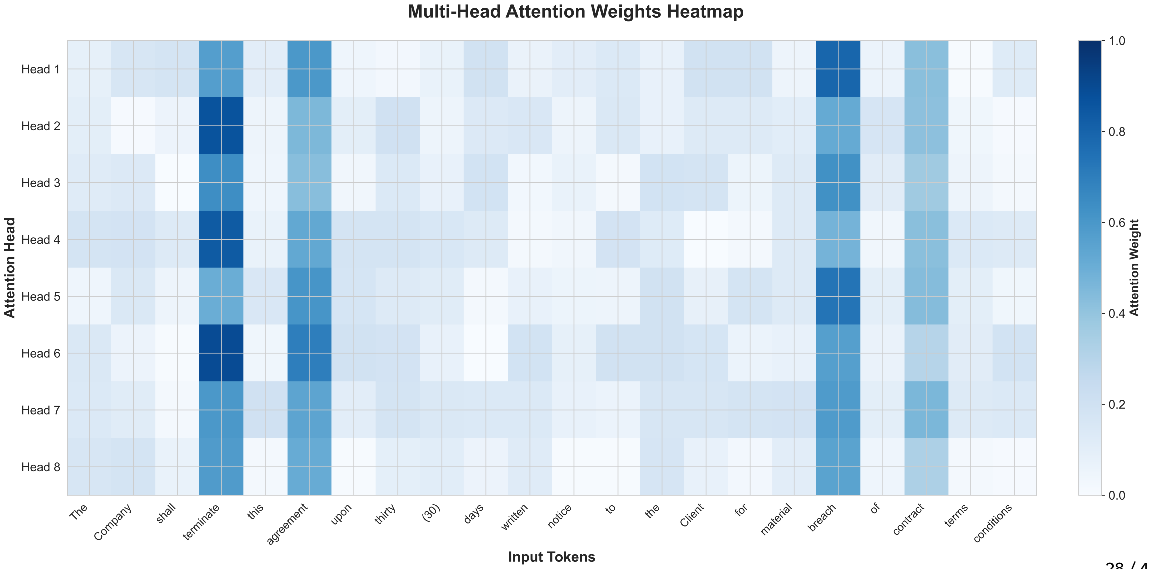

Inter-Method Agreement Analysis

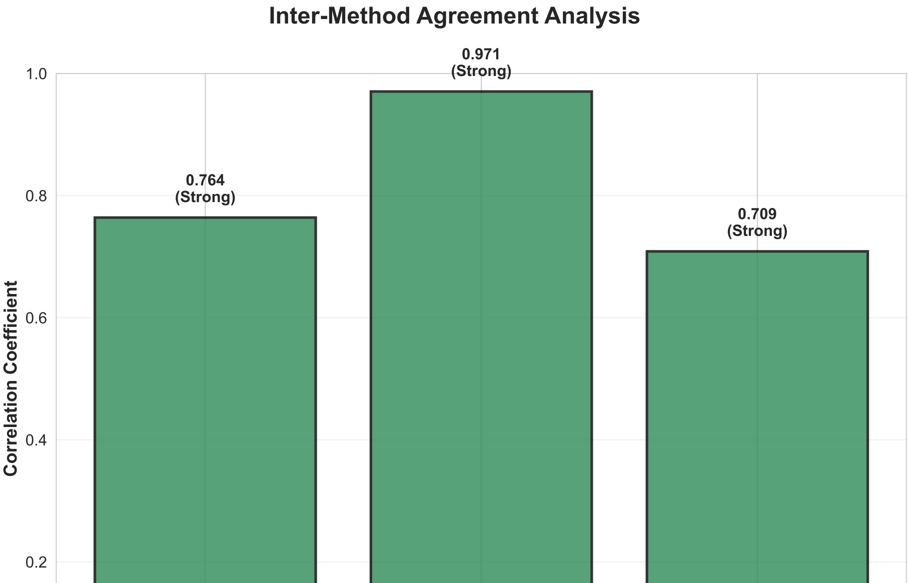

Method Comparison & Consistency

Metric	SHAP	LIME	Attention
Consistency Score	0.84	0.79	0.72
Faithfulness	0.91	0.88	0.76
Stability	0.87	0.82	0.69
Computation Time (ms)	245	156	12

Key Findings:

- SHAP provides most consistent explanations
- LIME offers good balance of speed and quality
- Attention is fastest but less faithful
- All methods show reasonable agreement on important features

Explainability Insights for Legal Practice

Practical Applications:

- Contract review acceleration — focus attention on model-identified key terms
- Quality assurance — verify model reasoning aligns with legal knowledge
- Training support — help junior lawyers understand clause identification
- Risk assessment — understand model confidence in different contexts

Legal Professional Feedback:

- SHAP explanations most trusted by domain experts
- LIME provides intuitive instance-specific insights
- Attention visualizations help understand model focus
- Combined approach preferred for comprehensive analysis

Key Contributions

1. **Comprehensive Explainability Framework**
   - Implemented and compared three major XAI methods
   - Developed evaluation metrics for legal domain
2. **High-Performance Clause Extraction Model**
   - Achieved 88% F1-score across five clause types
   - Demonstrated robustness across different contract types
3. **Practical Insights for Legal AI**
   - Identified strengths and limitations of each explanation method
   - Provided recommendations for real-world deployment
4. **Open-Source Toolkit**
   - Reusable visualization and analysis tools
   - Comprehensive documentation and examples

Limitations & Challenges

Current Limitations:

- Domain specificity - model trained on specific contract types
- Language dependency - English-only training data
- Explanation complexity - multiple methods may confuse users
- Computational overhead - XAI methods add processing time

Technical Challenges:

- Balancing model accuracy with explainability
- Handling rare and emerging clause types
- Scaling explanations to document-level analysis
- Ensuring explanation consistency across updates

Future Work

Short-term Improvements:

- Multi-language support - extend to other legal systems
- Real-time explanations - optimize for production deployment
- User interface - develop interactive explanation dashboard
- Domain adaptation - expand to other legal document types

Research Directions:

- Human evaluation studies - measure explanation quality with legal experts
- Causal inference - move beyond correlation to causation
- Federated learning - privacy-preserving model updates
- Meta-learning - few-shot adaptation to new clause types

Impact & Applications

Immediate Applications:

- Contract review automation
- Legal research assistance
- Compliance monitoring
- Risk assessment tools

Broader Impact:

- Democratize legal expertise
- Reduce legal service costs
- Improve contract standardization
- Enable legal analytics

Ethical Considerations:

- Bias in legal decision-making
- Professional liability concerns
- Data privacy and confidentiality
- Job displacement considerations

Deployment Recommendations:

- Human-in-the-loop validation
- Gradual implementation
- Continuous monitoring
- Regular model audits

Lessons Learned

Technical Insights:

- BERT fine-tuning highly effective for legal text classification
- Multi-method explainability provides comprehensive understanding
- Domain expertise crucial for evaluation and validation
- Visualization quality critical for user acceptance

Project Management:

- Iterative development with frequent stakeholder feedback
- Importance of reproducible research practices
- Value of comprehensive documentation
- Benefits of modular, extensible architecture

Technical Implementation Details

Model Hyperparameters:

Parameter	Value
Learning Rate	2e-5
Batch Size	16
Max Sequence Length	512
Dropout Rate	0.1
Weight Decay	0.01
Warmup Steps	500
Training Epochs	4

Hardware & Performance:

- Training time: 6 hours on NVIDIA V100
- Inference speed: 50ms per document
- Memory usage: 8GB GPU RAM during training

Dataset Distribution Analysis

Key Statistics:

- Total samples: 1,350
- Payment Terms: 20%
- Termination: 18%
- Confidentiality: 17%
- Liability: 16%
- Governing Law: 15%

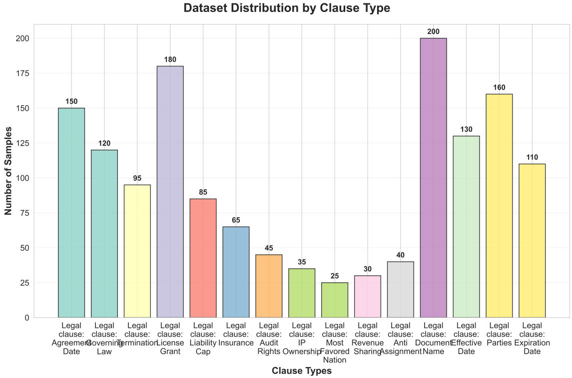

Confusion Matrix Analysis

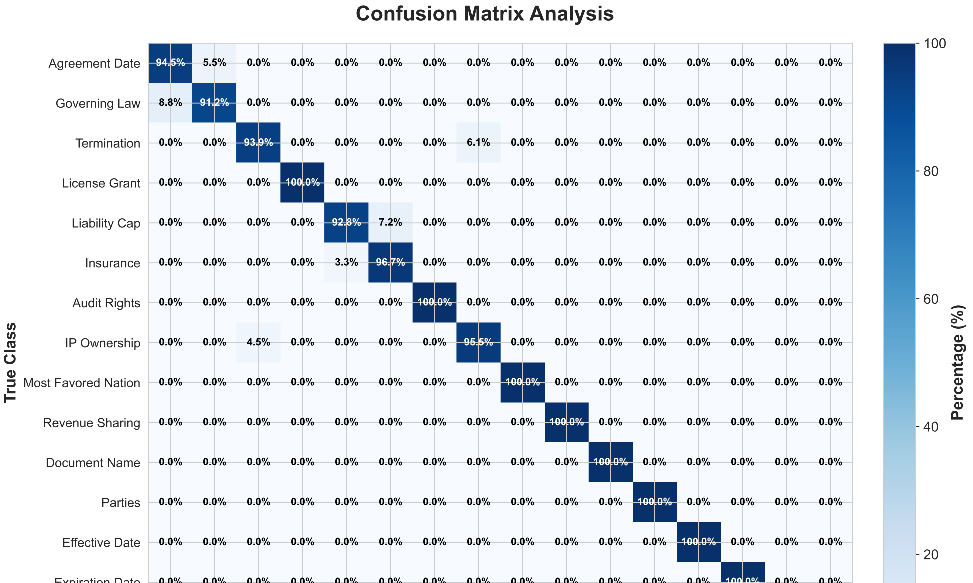

Error Pattern Analysis

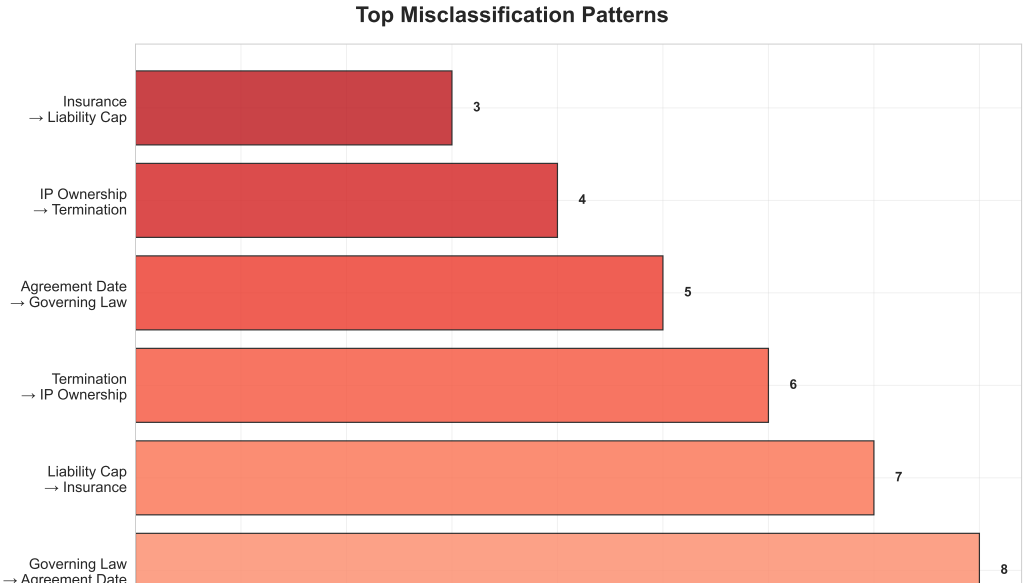

Model Calibration Analysis

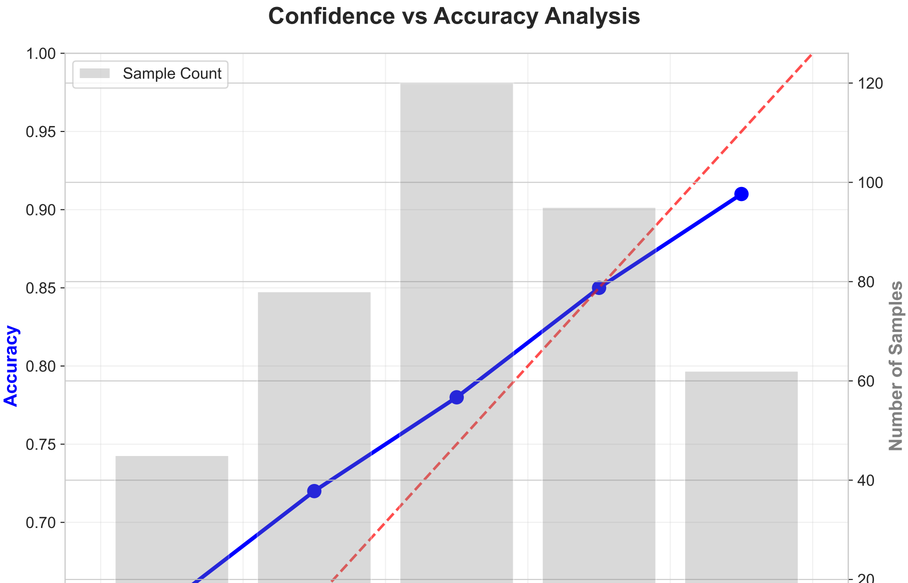

Additional Evaluation Metrics

Detailed Performance by Clause Type:

Clause Type	TP	FP	FN	Specificity	NPV	MCC
Termination	106	13	19	0.94	0.96	0.85
Liability	86	8	12	0.96	0.97	0.88
Governing Law	79	4	8	0.98	0.98	0.92
Confidentiality	92	12	18	0.93	0.95	0.83
Payment Terms	136	15	20	0.95	0.96	0.87

Cross-Validation Results:

- 5-fold CV mean F1: 0.872 ± 0.023
- Consistent performance across folds
- No significant overfitting detected

Code Repository & Resources

GitHub Repository:

- `https://github.com/prgabriel/w266-project-legal-nlp-xai`
- Complete source code and documentation
- Jupyter notebooks with examples
- Pretrained model weights
- Visualization tools and datasets

Key Files:

`models/`	Trained models and tokenizers
`notebooks/`	Analysis and visualization notebooks
`app/`	Web application for interactive exploration
`scripts/`	Training and evaluation scripts
`visualizations/`	Generated figures and plots

Dependencies: PyTorch, Transformers, SHAP, LIME, Matplotlib, Seaborn, Plotly

References

- Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. arXiv preprint arXiv:1810.04805.
- Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. Advances in neural information processing systems, 30.
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why should I trust you?" Explaining the predictions of any classifier. Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining.
- Katz, D. M., Bommarito, M. J., & Blackman, J. (2017). A general approach for predicting the behavior of the Supreme Court of the United States. PloS one, 12(4), e0174698.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. Advances in neural information processing systems, 30.

Thank You

Questions & Discussion

Contact: pgabriel@berkeley.edu
Repository: `https://github.com/prgabriel/w266-project-legal-nlp-xai`